

OC Core Methods and Reproducibility Companion v1.3.3

Methods, validation, reproducibility, and replay companion

Alexander Yashin

Independent Researcher

ORCID 0009-0008-6166-0914

Version 1.3.3

Manuscript revision date: 4 May 2026

Concept DOI: 10.5281/zenodo.17899134

Dedicated to my dear wife Maria, without whom this work would have been impossible.

Acknowledgements

The author thanks the reviewers and critics whose questions, objections, and suggestions contributed to the development of the Ontology of Continua and to the refinement of this release. The acknowledged contributors are G. V. Apostolov, Eduard Fadeev, Gennady Alekseevich Nosov, Sergey Shpadyrev, and Stanislav Tsukrov. Their criticism helped sharpen the claim boundaries, improve the reader route, expose weak presentation choices, and force clearer separation between model claims, proof routes, numerical evidence, prior-art comparison, and publication form. Acknowledgement records review pressure and intellectual contribution to the development of the work; it does not imply authorship, endorsement, publication approval, responsibility for the theory, or agreement with any claim promoted in the manuscript.

Abstract

This methods companion explains how the formal, finite, numerical, and archival checks support the promoted claims. Many contemporary systems are too entangled to be understood by a single domain vocabulary: a technical platform carries physical infrastructure, software architecture, institutional incentives, security boundaries, user cognition, and economic feedback at the same time. The Ontology of Continua addresses that situation as a typed model-core project. It describes continuants, realizations, liveness, death, residue, boundaries, morphisms, operators, cycles, dimension, and K-level witnesses under explicit assumptions, so that a reader can ask how a system persists, changes, fails, and reappears without changing languages at every disciplinary border.

The project is intentionally ambitious in subject matter and intentionally bounded in claim discipline. Its publication task is not to replace physics, biology, cognition, systems engineering, enterprise architecture, or economics with a slogan. Its task is to offer a common structural grammar through which domain experts can compare state spaces, thresholds, flows, cycles, boundaries, and failure conditions without pretending that domain-specific science has become unnecessary.

Version 1.3.3 is a bounded external-review release. It is mature enough to be read as a coherent scientific manuscript and to be checked against formal and executable artifacts, but it remains a staged research program rather than a declaration of final all-domain completion. The release promotes model-core claims where their assumptions and evidence are visible; it leaves broader numerical closure, wider domain validation, and stronger comparative claims as future obligations.

The package contains a full monograph, a compact article route, a methods and reproducibility companion, a reviewer attack-and-response map, release navigation, public evidence anchors, finite-model outputs, Lean subset evidence, bounded target-blind replay QA, comparator and prior-art material, and appendices for proof, figures, tables, numeric rows, and source trace. The monograph is the canonical reading spine; the other documents are curated routes for first reading, editorial review, replay, and hostile criticism.

The evidence status is deliberately mixed and explicit. Some claims are supported by theorem routes, proof sheets, Lean declarations, and finite semantic witnesses. Some claims are supported by bounded numerical or target-blind replay rows that should be read as scoped QA evidence, not as complete empirical validation of a whole domain. Some claims are positioned against prior art or retained as research boundaries because the available evidence is not yet strong enough for a broader statement.

The principal limitation is therefore part of the manuscript's method: public wording must not

outrun the evidence. OC Core 1.3.3 asks the reader to evaluate whether the typed model, proof governance, finite semantics, numerical anchors, figures, tables, appendices, and reviewer-response routes form a coherent bounded model-core contribution. Future work must add evidence or mechanization before it widens the public claim surface.

The practical value of the manuscript is therefore a form of disciplined compression. If the model is useful, it should help readers translate complex systems into a shared ontology while preserving the meaningful distinctions that make each domain scientifically serious. That is the standard by which the release asks to be read.

Version 1.3.3 Release Delta

The OC Core release sequence is part of a continuing scientific program rather than a series of isolated archive events. As the model is tested against formal criticism, domain examples, reproducibility checks, and external review, the public manuscript is expected to become more precise, more explicit about its limits, and more useful to readers outside the author’s immediate working context.

A new release is therefore warranted only when there is a meaningful scientific or editorial delta: a clarified definition, stronger proof route, better evidence boundary, improved reader pathway, repaired notation, new domain projection, or more honest comparator position. Minor process movement is not enough. The public release should tell readers what changed in the scientific object and why that change matters.

This policy is also a commitment to regularity without pretending that research can be made mechanical. The author intends to make future public updates systematic enough for readers, reviewers, and auditors to follow the development of the model, while preserving the difference between a public scientific result and private working material. No internal route, unpublished process detail, or control-plane record is promoted as reader-facing evidence merely because it helped produce the release.

Version 1.3.3 is the release in which the OC Core corpus is reorganized from a set of scattered source witnesses and process records into a reviewable scientific package. The visible delta is not merely a new archive number: the release strengthens the typed foundation, makes the claim boundary explicit, integrates the theorem route with public proof sheets, and separates reader-facing prose from machine evidence.

The model delta includes clearer treatment of carriers, realizations, liveness, death, residue, re-birth, morphisms, generalized boundaries, hybrid operators, cycle modes, historical and effective dimension, and K-level witnesses. These additions are presented as a bounded model-core grammar rather than as an unrestricted theory of every phenomenon.

The verification delta adds and consolidates a Lean-checked subset, structured proof sheets, finite semantic checks, finite witness interpretation, bounded target-blind replay QA, comparator and prior-art positioning, negative-control and falsifier language, and adversarial-review response material. Where a stronger claim is not yet earned, the release records the limitation instead of hiding it inside a source-control record.

The editorial delta is equally important. Figures, tables, formulas, numeric anchors, appendices, and evidence routes are kept in the publication corpus; machine coverage maps remain coverage and trace data only. The public documents are therefore built from the human manuscript hierarchy

and curated payload sources, while source bindings and machine evidence indexes stay in the review manifest.

For a returning reader, the practical point is this: 1.3.3 should be read as a clarification and consolidation release. It does not claim final closure. It improves the public surface through which the model can be criticized, taught, checked, and extended.

Reader Routes

Dear readers, systems theory is of interest to many communities, but rarely in exactly the same way. A formal reviewer, a philosopher of systems, an applied architect, an AI engineer, a security specialist, and an executive reader may all ask legitimate questions of the same manuscript. OC Core 1.3.3 is therefore designed as a deliberately generous scientific reading surface: it aims to disclose the Ontology of Continua as an applied model of system architecture while giving different readers a courteous route into the material.

Scientific reviewers and formal critics may wish to start with the abstract, release delta, model-foundation route, theorem/proof integration route, Lean subset, finite semantic witnesses, proof machinery, and falsifiability sections. This route is meant to make the work attackable in the best sense: every promoted claim should expose its assumptions, proof or executable evidence, counterexample boundary, and reopening condition.

Systems theorists, philosophers, and interested theoretical readers may prefer to start with the continuum problem, the historical and systems-theory motivation, the K-level primer, lifecycle and boundary chapters, prior-art comparison, and limitations. This path asks what OC inherits, what it reorganizes, where its residual delta begins, and where predecessors or parallel branches already carry part of the work.

Applied architects of complex systems—including engineers, enterprise architects, AI builders, security specialists, and applied researchers—may find it useful to start from practical examples, domain projections, figures, tables, and worked routes before returning to the formal core. The model chapters can then be read as a design grammar for carriers, realizations, boundaries, operators, update/flow separation, evidence classes, and claim boundaries.

CIOs, CEOs, enterprise architects, and strategy readers may sensibly read the release delta, practical utility chapter, model-comparison appendix, and final conclusion before investing in proof details. This route asks what the model is for, what it can support today, what remains too immature for operational commitment, and how evidence classes affect research or enterprise decisions.

Reproducibility auditors, evidence reviewers, and benchmark readers may go directly to the methods and evidence chapters, target-blind replay QA, numeric tables, machine-readable table reference, audit trail, source-trace appendix, and public evidence package manifest. This route asks whether formula, input snapshot, comparator, residual or uncertainty, negative control, falsifier, replay hash, and failure interpretation are present.

The companion is useful when commands, inputs, outputs, checksums, finite witnesses, target-blind replay, and failure interpretation are the main concern. In the full package, the conceptual model and formulas live in the monograph’s scientific body; the proof route and Lean subset live in theorem, proof, and formalization sections; numeric evidence and target-blind replay QA live in the methods/evidence route; figures and tables live inline in the monograph; appendices carry

source trace, proof machinery, comparison rows, audit trail, and machine-readable table references; prior-art comparison and reviewer objections live in the article and attack-response map.

The author asks all readers to treat limitations as part of the claim, not as a footnote. A statement is promoted only where its assumptions, proof or evidence class, comparator boundary, and reopening condition are visible. Claims about complete scientific coverage, unrestricted all-domain numerical closure, or unrestricted superiority over contemporary science are not promoted by this release.

Table of Contents

Acknowledgements	1
Abstract	1
Version 1.3.3 Release Delta	2
Reader Routes	3
1 Methods and Reproducibility Companion	6
1.1 What a Replay Checks	6
1.2 Formal and Finite Checks	6
1.3 Interpreting Failure	6
1.4 Audit Trail and Machine-Readable Tables	6
1.5 Governed Local Review Boundary	6
1.6 Local Editorial Review Status	6
1.7 Journal-Facing Projection Boundary	7
1.8 Visual Argument Boundary	7
1.9 Table and Quantitative Surface Boundary	7

1 Methods and Reproducibility Companion

The methods companion explains how a reader can inspect OC Core 1.3.3 without treating machine-readable material as the manuscript itself. Reproducibility is used here as scientific interpretation: a check matters because it bears on a named claim, assumption, formula, witness, comparator, or falsifier.

1.1 What a Replay Checks

A replay row connects an input snapshot, a formula or reconstruction rule, an expected output, a residual or uncertainty, a comparator, a negative control, and a failure interpretation. Passing such a row supports the scoped claim named by the row. It does not by itself validate an entire domain or remove the need for broader empirical work.

1.2 Formal and Finite Checks

The proof route and finite semantic checks give the model a stricter ceiling. A theorem statement is supported only when the assumptions, proof sheet, dependency relation, mechanized subset where present, and counterexample boundary remain inspectable. A finite witness is valuable because it can show that a definition or operator behaves as claimed inside a bounded semantics.

1.3 Interpreting Failure

A failed check is not a formatting problem. It may reopen a claim, narrow a domain route, reveal a missing assumption, or show that a comparator already explains the case. The companion therefore treats mismatch, missing input, missing hash, drifted output, absent negative control, and unsupported extrapolation as scientific events.

1.4 Audit Trail and Machine-Readable Tables

The reader-facing tables summarize the evidence in prose. Machine-readable tables remain available for auditors who need to inspect source rows directly: claim rows, theorem rows, finite checks, target-blind replay QA, numeric tables, comparator material, and source-trace material. The tables are support for the argument, not a substitute for it.

1.5 Governed Local Review Boundary

This reader-facing text is not a transcript of local model output. A governed local review service is used as a bounded review and critique surface for bounded chunks after deterministic checks. Its role is to help detect leakage, weak reader prose, and missing anchors under host-safety policy; the public claim remains the manuscript’s responsibility.

1.6 Local Editorial Review Status

The reader-facing package has been reviewed through a governed local editorial route in small packets before higher-level aggregation. That route is used to detect leakage, weak didactic prose, missing anchors, and reader-facing defects while preserving the manuscript as the public scientific surface.

1.7 Journal-Facing Projection Boundary

This artifact is part of the OC Core release SPOT: the common source surface from which venue-specific reader packages are projected. Explanatory prose may be adapted for a venue, but the scientific claim, formula, figure, table, proof, evidence, and citation anchors remain bound to the release corpus. The venue package is prepared for owner review only and performs no external action.

1.8 Visual Argument Boundary

Figures in this package are treated as part of the scientific argument. A figure must separate its labels, identify the variables or numbers it makes visible, and point back to the formula, proof, table, or evidence route it supports. The manuscript therefore uses figures as compact explanations rather than as decoration.

1.9 Table and Quantitative Surface Boundary

Tables in this package are treated as argument surfaces rather than raw data dumps. A promoted reader-facing table must have a readable layout, a captioned scientific role, and a source or evidence anchor; internal probe files used to test table geometry remain outside the publication package.

Keywords and Citation Route

Keywords: Ontology of Continua; continuum ontology; typed model core; systems theory; autopoiesis; dynamical systems; hybrid systems; formal methods; Lean formalization; finite semantic checks; target-blind replay QA; reproducible research; artifact evaluation; claim governance; falsifiability; evidence-bound scientific publishing.

Citation identity: Alexander Yashin, OC Core Methods and Reproducibility Companion v1.3.3, 4 May 2026. The Concept DOI is printed on the title page.

Open repository: <https://github.com/alexanderyashin/ontology-of-continua-core-main>. Repository files support reproducibility and source inspection; they do not replace the manuscript's claim boundaries.